# A Pre-Registered Cost Model for a Small Bit-Exact int8 CNN Accelerator on a RISC-V Microcontroller SoC

Anirudh Sahu
anirudh12sahu@gmail.com

***Abstract*—We present MCUACCEL, an int8 convolution, depthwise and fully connected accelerator memory-mapped into a PicoRV32 SoC, together with a backend for `mcugrad` that emits accelerator command descriptors and driver C while keeping every other operator in the original software kernels. The accelerator is bit-exact with the TFLite-style integer reference. Its contribution is methodological: an analytic port-bound cycle model, with per-command tolerances, was committed to git (tag `m1`, 6533b1e) before any datapath RTL existed, and the RTL was then scored against it. At $P = 8$ output channels per tile the frozen model predicts a speedup of 27.0× on `mnist_cnn` and 363.1× on `cifar_mobilenet` against the software baselines measured at the time; after a later fairness fix to the build (decision D19) the baselines are 11,383,383 and 69,169,620 cycles per inference and the measured speedups are 46.8× and 350.0×. A post-hoc rescoring that substitutes only the measured post-D19 fallback-operator cycles into the frozen prediction gives 46.7× and 348.0×. Every accelerator command across five values of $P$ is predicted within 12.4 % (median signed error 0.2 %), and the per-network accelerator-plus-driver time within -0.9 % and 0.4 %. Claim C1 is confirmed; C4a is confirmed, although 6 commands exceeded their pre-registered bias (root cause given post hoc); C3 is refuted: 7,313 LUT4 against a budget of 6,000. All cycle counts are Verilator cycles, not FPGA wall-clock; in our ECP5 flow the SoC closes timing at 35.8 MHz with the accelerator and 61.0 MHz without it, a factor any wall-clock speedup must absorb.**

## I. Introduction

Quantized CNNs cost tens of millions of cycles per image on microcontroller-class cores. That a dedicated int8 accelerator beats a scalar in-order RISC-V core by an order of magnitude or more is well established [1]–[3]; reproducing that is not our contribution. We ask instead how well a simple analytic model, written down *before* the RTL, predicts the cycle counts of a small memory-bound accelerator, with tolerances fixed in advance, and we report the outcome whether or not it flatters the design. We found no prior accelerator evaluation that froze its cycle model and error tolerances before the RTL existed (Sec. II); the model itself is not claimed as novel. The claims below are pre-registered in `docs/SPEC.md`; the predictions are frozen at tag `m1` (commit `6533b1e`, `m1_predictions.csv` sha256 prefix `fd94f39a1a913ee6`).

- C1 Speedup at $P = 8$ is at least 10× on both networks, end to end with fallback layers and driver. **C1b**: at least 100× on `cifar_mobilenet`.
- C2 Bit-exactness: zero mismatching bytes on 1,000 random layers (C2a), in every intermediate tensor of both networks (C2b), and in final outputs for 100 images per network on the full SoC (C2c).
- C3 At $P = 8$ the accelerator uses at most 6,000 LUT4, 40 `MULT18X18D` and 12 `DP16KD` on an ECP5 LFE5U-85F.
- C4 The frozen model predicts every accelerator command within 20 % (C4a, all $P$) and each network's accelerator-plus-driver time within 10 % (C4b).
- C5 The saturation width of the $P$ sweep is reported, predicted versus measured; its mechanism was stated in advance.

**Contributions.** (i) A from-scratch int8 accelerator (conv, depthwise, fully connected) with chained in-memory descriptors, bit-exact with the gemmlowp rounding rule, and a backend that compiles `mcugrad` graphs to descriptor chains plus driver C with software fallback for everything else (Sec. III, IV). (ii) A closed-form cycle model of the accelerator, its driver and the fallback operators (Sec. V), frozen before the datapath, with a written list of its known biases. (iii) A measurement protocol in which deviations are logged rather than absorbed, including a build-fairness fix that moved the baseline after the predictions were frozen (Sec. VI). (iv) The scored outcome of every claim, the refuted one included, with a post-hoc mechanism for each of the 6 commands that exceeded its registered bias (Sec. VII). Every number in this paper is generated from result files by a script in the repository.

## II. Background and Related Work

**Quantization.** The numerical contract is the integer-only scheme of Jacob et al. [4] as implemented in gemmlowp [5] and specified by LiteRT [6]: a saturating rounding doubling high multiply followed by a rounding divide by a power of two. **MCU inference.** TFLite Micro [7], CMSIS-NN [8] and MCUNet [9] provide software kernels; `mcugrad` [10] is the compiler whose int8 simulator is our specification. **Accelerators.** Eyeriss [11], the TPU [12], NVDLA [13], VTA [3], Gemmini [1] and CFU Playground [2] show large speedups over in-order cores; FINN [14] and PULP-NN [15] target FPGA and multicore MCU settings. Closer in scale, XpulpNN [16] extends the RISC-V ISA instead of adding a bus master, and a depthwise-separable accelerator for a RISC-V core [17] (preprint) targets the layer type that dominates

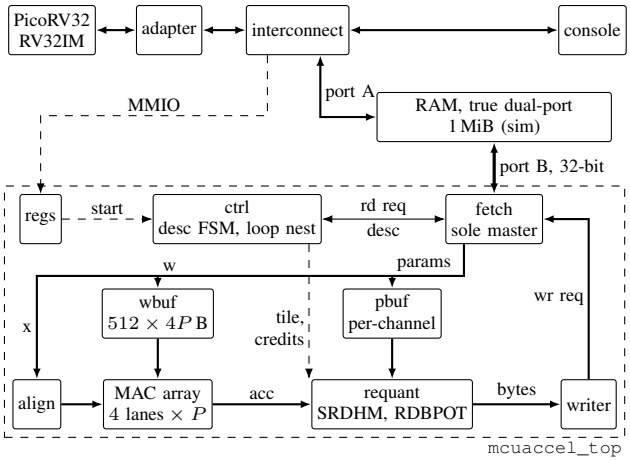



Fig. 1. System and accelerator block diagram. Boxes are RTL modules (vendored PicoRV32 and `soc_*` above, `mcuaccel_*` in the dashed box); solid arrows carry data, dashed ones commands. The CPU writes a descriptor address and START through MMIO, then polls; the accelerator reaches RAM only through its private port B.

`cifar_mobilenet`. These works report the speedup of a new design; speedup is not our contribution. None of their RTL is used here. **Models.** Earlier analytic models [18], [19] were developed together with their designs and validated against the implementation, and the roofline model [20] bounds rather than predicts; ours was frozen first, which is the practice of pre-registration transplanted to hardware evaluation. The CPU is PicoRV32 [21] on a Lattice ECP5 [22], with the bus protocol of corv32 [23].

## III. Architecture

The CPU reaches a dual-port RAM through an adapter and interconnect; the accelerator is a second bus master on the RAM's other port (a private 32-bit port, one word per cycle, one outstanding request), so no arbiter is needed. The CPU writes a descriptor address and a start bit, then polls a status register (Fig. 1).

**Dataflow.** The array is output-stationary over one output pixel and a tile of $P$ output channels. In NHWC layout the in-bounds part of a kernel row is contiguous, so each cycle one four-byte activation word meets one weight-buffer row of $4 \cdot P$ bytes and performs $4 \cdot P$ multiply-accumulates; every channel count uses one mechanism, and an aligner absorbs the three non-zero word alignments at the cost of one priming read per kernel row. Depthwise layers get only 4 useful MACs per cycle. Weights and per-channel parameters are loaded once per command into local buffers (512 rows). A single requantizer emits one byte per cycle using the specification's SRDHM/RDBPOT rule; accumulators are 32 bit and the backend enforces a no-wrap bound. Two credits separate the sequencer from the requantizer: a tile's last request takes a credit, and the requantizer returns it when the tile's last value enters its pipeline.

**Interface.** A command is a 18-word descriptor in RAM (geometry, strides, addresses, zero points, clamp). Descriptors chain by pointer and the accelerator writes a cycle stamp into each, so per-command time is observable without a trace. Anything outside the conv/depthwise/fully-connected envelope falls back to software.

## IV. Backend and Driver

The backend (`backend/mcuaccel_backend`) consumes mcugrad's int8 graph, packs weights and parameters, splits layers that exceed buffer limits, builds descriptors, and emits a `model.c` with the same `model_run` API; unsupported ops are the identical mcugrad kernels. A pure-Python executor of the descriptor semantics, tested equal to the integer reference, is the oracle for the RTL's request stream. The driver writes the descriptor address, starts the chain, polls, and patches only the fields that depend on the network input and output. Its cost, measured on the walking-skeleton stub before the model was frozen, is 71 cycles per call plus 15 per patched field.

## V. Cost Model

The model (`model/m1.py`; normative text in `docs/DESIGN.md` Sec. 7) is a loop nest over the output pixels and channel tiles of a command, with one closed form per tile. Let $\lambda$ be the cycles per bus transaction ($\lambda = 1$ on the zero-wait RAM), $R$ the number of activation words a pixel needs for one tile (in-bounds kernel rows, plus a priming word per misaligned row), $W_t$ the output words written for tile $t$, $NV_t$ the number of valid outputs of the tile ($\min(P, \mathrm{NCH} - tP)$ for conv, 4 for depthwise), $D = 18$ the descriptor length, $W_\mathrm{w}$ the weight words loaded, and NCH the command's output channels. Then

$$T_\mathrm{tile}(t) = \max\big(\lambda(R{+}W_t),\, NV_t,\, \lceil \tfrac{NV_t+\lambda+3}{2} \rceil\big), \tag{1}$$

$$T_\mathrm{pix} = 1 + \textstyle\sum_t T_\mathrm{tile}(t), \tag{2}$$

$$T_\mathrm{cmd} = \big[(D{-}1)\lambda + 1\big] + \big[\lambda W_\mathrm{w} + 1\big] + \big[3\lambda\,\mathrm{NCH} + 1\big] + \textstyle\sum_\mathrm{pix} T_\mathrm{pix} + \big[2\lambda + 10 + NV_\mathrm{last}\big], \tag{3}$$

$$T_\mathrm{chain} = \textstyle\sum_\mathrm{cmd}(T_\mathrm{cmd} + \lambda + 1) + 2. \tag{4}$$

The three terms of (1) are the port bound (reads of this tile plus the writes of the previous one, which steal port slots), the requantizer bound (one output per cycle) and the two-credit latency bound for tiles with few outputs, where $\lambda + 3$ is the issue-to-accumulator latency. The 1 in (2) is the per-pixel setup cycle. The bracketed terms of (3) are the descriptor, weight and parameter loads and the final drain, in order; $T_\mathrm{cmd}$ is what the accelerator writes as the command's stamp. Geometry enters only through $R$, $W_t$ and $NV_t$, so the whole model fits in a page of the design document and was recomputed by hand for three layers by a reviewer before freezing. For a $1 \times 1$ convolution with aligned channels, (1) reduces to $\max(C/4 + \lceil P/4 \rceil,\, P,\, \cdot)$: the design is port bound when $C$ exceeds about $3P$ and requantizer bound below that, which already predicts that $P$ saturates early on networks with narrow layers.

**Whole network.** The predicted accelerated `model_run` is the sum over chains of $T_\mathrm{chain}$ plus 71 plus 15 per patched field, plus the measured baseline cycles of every operator that stays

in software, plus the 120-cycle bench window. The software term is a measurement carried through for the speedup figure; claim C4b scores only the accelerator-plus-driver part, so that the prediction is not trivially right because its largest term was measured.

**Registered biases.** Before freezing, the design document listed where the model is known to be off and by how much: the setup cycle is hidden behind the requantizer on requantizer-bound layers (an over-prediction of up to a few percent on named commands and under 0.5 % on all others); the drain term double counts the last tile's writes (at most $P + 3$ cycles per command); the driver constants omit a bind check and phase wait of a few tens of cycles; and the measured interval contains profiling macros that the prediction does not. Each is a stated direction and magnitude, so a measured error can be checked against it, which Sec. VII does.

## VI. Methodology

**Pre-registration.** Tag `m1` precedes every commit that adds a datapath or sequencer. The model is not changed to match the RTL, nor the RTL tuned to the model: the RTL is scored at its first full measurement and later cycle-changing edits are logged with before/after values. Predicted rows are those of `m1_predictions.csv` with architecture `os` (17 commands at $P = 8$, 108 over the sweep). The tag and the file hash were made public at https://github.com/anirudh12032008/mcuaccel/tree/m1 on 2026-11-01; that timestamp shows the predictions existed by then, while their order before the datapath rests on the git history.

**Timing protocol.** Inference time is the `rdcycle` interval around `model_run` on the SoC, including software layers, descriptor patching, polling and driver cost, over 10 fixed test images and 3 timed runs after a warm-up (decision D17: the simulator is deterministic, so more runs add nothing). We report the median of per-run per-image medians with min and max. The bench window costs 120 cycles in both builds. Cycle counts come from Verilator simulation of the full SoC at a nominal 50 MHz. PicoRV32 is configured with `ENABLE_FAST_MUL`, `ENABLE_DIV` and `BARREL_SHIFTER`; firmware is built with GCC 15.2 at `-O2 -march=rv32im_zicsr` (both recorded in each result file).

**Deviations after m1.** Two logged changes altered measured cycles; neither changed the frozen predictions or a threshold (values in `results/deviations.json`, written by `experiments/make_deviations.py` from the decision text with the after-values checked against `results/`). *D19:* identical fallback source compiled to different code in the two builds, because inlining depended on which operators a file held. After pinning (below) the baseline moved from 12,188,028 to 11,383,383 (mnist) and 72,157,425 to 69,169,620 (cifar) cycles, and the accelerated $P = 8$ build from 451,480 to 243,059 and 197,617 to 197,613; measured speedup moved from 27.0× to 46.8× (mnist) and 365.1× to 350.0× (cifar). The mnist change is large because its accelerated runtime is dominated by the pooling and transpose fallbacks, which the old inlining compiled badly; no accelerator cycle changed. *D20:* the RTL's credit return had one register stage too many against the normative timing (Sec. 7.2 of the design); it was made combinational. Smoke runs on the earlier RTL gave a C4b error of -1.0 % (mnist) and +0.4 % (cifar); the scored RTL gives -0.9 % and +0.4 %.

TABLE I
Software baseline on the SoC (cycles per inference).

| Network | median | min | max |
|---|---|---|---|
| `mnist_cnn` | 11,383,383 | 11,378,852 | 11,387,883 |
| `cifar_mobilenet` | 69,169,620 | 69,144,241 | 69,186,835 |

**Fairness.** The baseline is mcugrad's generated `model.c` with two build-time attribute changes (decision D19): the `mg_*` kernels are `always_inline` and each `op_*` function is `noinline`, so a fallback operator is compiled from its own text alone; kernel and operator source text is unchanged. The same pinning is applied to the baseline and to the accelerated build, which use the same compiler and flags on the same SoC (accelerator present and idle in the baseline). Fallback-operator identity is checked by comparing per-op profile cycles of the two builds, not by object-code hash. Optional backend passes are off in every scored build.

## VII. Results

**Baseline.** Table I gives the software baseline. Convolution and fully connected layers account for 98.0 % (mnist) and 99.9 % (cifar) of baseline cycles in the profile build, which bounds the Amdahl ceiling.

**C1/C1b: speedup.** Table II compares the frozen prediction with measurement. At $P = 8$ the prediction puts the software ops at 95.8 % (mnist) and 31.6 % (cifar) of the accelerated run: pooling and layout ops that stay in software dominate MNIST, whose predicted speedup saturates near 27.2× at the largest $P$. C1 is confirmed; C1b is confirmed. The frozen predictions in the table were computed against the pre-D19 baselines (12,188,028 and 72,157,425 cycles) and pre-D19 fallback-operator cycles, so they are not comparable one to one with the measured columns; the last row is a post-hoc rescoring (frozen accelerator-plus-driver prediction, plus the measured fallback-operator cycles of the accelerated build, plus the 120-cycle bench window): 243,815 and 198,760 cycles against measured 243,059 and 197,613.

**C2: bit-exactness.** Random layers: 5,000 run, 0 mismatching bytes (C2a confirmed). Intermediate tensors: 10 images, 0 mismatches. Full SoC: 100 images, 0 mismatches against the integer reference and 0 (mnist), 0 (cifar) against the baseline firmware's outputs (C2c confirmed).

**C3: resources.** Table III reports the Yosys/nextpnr-ecp5 flow on the LFE5U-85F. The accelerator at $P = 8$ uses 7,313 LUT4, 36 `MULT18X18D` and 11 `DP16KD`; the multiplier and block-RAM budgets hold, the LUT budget does not, so C3 is refuted. The LUT budget followed a pre-RTL estimate in the design document; we report the miss rather than revise the

TABLE II
SPEEDUP VERSUS $P$: FROZEN PREDICTION (AGAINST PRE-D19 BASELINES) AND MEASUREMENT; LAST ROW: POST-HOC RESCORING AT $P = 8$ IN THE PREDICTION COLUMNS.

| | mnist | | cifar | |
|---|---|---|---|---|
| $P$ | pred. | meas. | pred. | meas. |
| 1 | 22.9 | 35.2 | 111.7 | 107.2 |
| 2 | 25.2 | 41.2 | 187.1 | 179.3 |
| 4 | 26.4 | 45.0 | 282.0 | 271.2 |
| 8 | 27.0 | 46.8 | 363.1 | 350.0 |
| 16 | 27.2 | 47.5 | 411.4 | 398.4 |
| 8 post-hoc | 46.7 | 46.8 | 348.0 | 350.0 |

TABLE III
SYNTHESIS (RESOURCES FROM THE LOWEST SEED OF THE NEWEST RTL VERSION OF EACH DESIGN; FMAX MEDIAN OVER ITS DISTINCT SEEDS).

| Design | LUT4 | MULT | BRAM | Fmax MHz | seeds |
|---|---|---|---|---|---|
| PicoRV32 | 2,896 | 4 | 0 | 81.7 | 10 |
| Accel., $P = 1$ | 5,788 | 8 | 4 | 61.1 | 3 |
| Accel., $P = 8$ | 7,313 | 36 | 11 | 61.5 | 10 |
| Accel., $P = 16$ | 9,716 | 68 | 19 | 58.0 | 3 |
| SoC, no accel. | 3,392 | 4 | 128 | 61.0 | 10 |
| SoC with accel. | 10,684 | 40 | 139 | 35.8 | 10 |

budget. Across the sweep the LUT count grows from 5,788 at $P = 1$ to 9,716 at $P = 16$: the fetch, sequencer and requantizer are a fixed cost that the multiplier array only overtakes at the largest $P$. A post-hoc per-module synthesis attributes 45 % of the $P = 8$ LUT-equivalents to the sequencer (3,457 against a design estimate of 1,300), independent of $P$, while the MAC array grows by about 194 per lane (300 placed); the per-module counts are an attribution, since the hierarchy adds 971 over the flattened design and placement 518. Narrowing the sequencer's 32-bit address arithmetic to the SoC address width and keeping the requantizer's partial-product sum and product registers in DSPs are likely, unmeasured ways under the budget (post hoc).

**C4/C5: prediction accuracy.** Table V lists predicted command cycles at $P = 8$ against the stamp measurement. Over the commands of the sweep the maximum absolute error is 12.4 %, the median signed error 0.2 %, with 0 tolerance violations out of 108 rows (C4a confirmed). Per network, accelerator plus driver, the signed error is -0.9 % (mnist) and 0.4 % (cifar). Figure 2 overlays the measurement on the frozen prediction: summed STAMP per inference versus $P$, and every scored command against the $\pm 20\,\%$ band. The predicted saturation width is 16 (mnist) and 16 (cifar) on the grid, where 16 means not reached, against measured 16 and 16. C5 is descriptive; its falsifiable part is C4a at every $P$. The model attributes the flattening to the depthwise layers, the single requantizer and the weight load, not to the port.

**C4 post-hoc notes.** The pre-registered bias list (design Sec. 7.4) gave an over-prediction for a few commands and under 0.5 % for all others. 6 scored rows exceed their registered value by more than 0.5 percentage points (Table IV); all

TABLE IV
COMMANDS WHOSE ERROR EXCEEDS THE PRE-REGISTERED BIAS (CYCLES; ERROR AND REGISTERED BIAS IN %).

| net | op | $P$ | pred. | meas. | err. % | reg. % |
|---|---|---|---|---|---|---|
| mnist | 1 | 8 | 7,144 | 6,356 | 12.4 | 3.4 |
| mnist | 1 | 16 | 7,168 | 6,380 | 12.4 | 3.4 |
| mnist | 3 | 16 | 4,580 | 4,515 | 1.4 | 0.5 |
| cifar | 0 | 4 | 50,092 | 49,381 | 1.4 | 0.5 |
| cifar | 0 | 8 | 27,724 | 27,029 | 2.6 | 0.5 |
| cifar | 0 | 16 | 17,648 | 16,616 | 6.2 | 1.9 |

stay inside the tolerance, so C4a is unchanged. Root cause, found after measurement (post hoc, RTL trace of all 6 rows): in every row the deviation splits cycle for cycle into writes granted in the per-pixel setup cycle, requantizer-bound pixels and the drain. A pending write of the previous pixel granted in the setup cycle hides that cycle, so the RTL behaves as a maximum over $1 + \lambda R$, $\lambda(R + W)$ and the requantizer term, whereas (2) adds the setup cycle to the maximum; the registered bias counted only strictly requantizer-bound pixels. A tie of $\lambda(R+W)$ with the requantizer term is a special case and fully explains mnist op 1 at $P = 8$ and $P = 16$ (546 and 546 of 784 setup cycles carry a write), cifar op 0 at $P = 16$ (690 of 1,024) and mnist op 3 at $P = 16$ (24 tie writes among 47 of 196). Cifar op 0 at $P = 4$ and $P = 8$ is strictly port-bound (0 tie writes), yet 707 and 690 of 1,024 setup cycles carry a write. With all setup-cycle writes removed, no row registered at the default bias keeps more than 0.40 % error. The model outputs stay as scored (decision D21).

The per-chain cycle counter obeys $\text{CYCLES} = \sum(\text{STAMP} + \lambda + 1) + 1$ on 275 of 275 chain calls; the frozen check, with $+2$ as in (4), holds on 0. Root cause: the accepting cycle of the start write is not busy. RTL and STAMP are correct; the design's chain time with $+2$ includes that cycle and is the CPU-visible term, so C4b is unaffected beyond one cycle per chain call. An independent bus-level count, from the first descriptor request to the stamp request, equals the self-reported STAMP on 3 of 3 commands (6,356, 7,715, 16,616 cycles), the first such check.

## VIII. DISCUSSION AND LIMITATIONS

All cycle counts come from Verilator simulation; the FPGA is used only for synthesis figures (LUT, DSP, BRAM, Fmax). No FPGA wall-clock or power measurement exists yet; the procedure is in `HARDWARE_TODO.md`. Wall-time speedup would be the cycle speedup scaled by the Fmax ratio of the SoC with and without the accelerator, and Table III shows that ratio is below one: the SoC with the accelerator closes at 35.8 MHz against 61.0 MHz without it in this flow, and below the nominal 50 MHz, so a board result would be smaller than the cycle speedup by that factor. The memory model is zero-wait on-chip RAM; an external memory would change the port-bound assumption, although $\lambda$ in the model is the knob that would absorb it. The accelerator covers convolution, depthwise (multiplier one) and fully connected layers only; pooling and layout ops stay in software and cap MNIST by Amdahl's

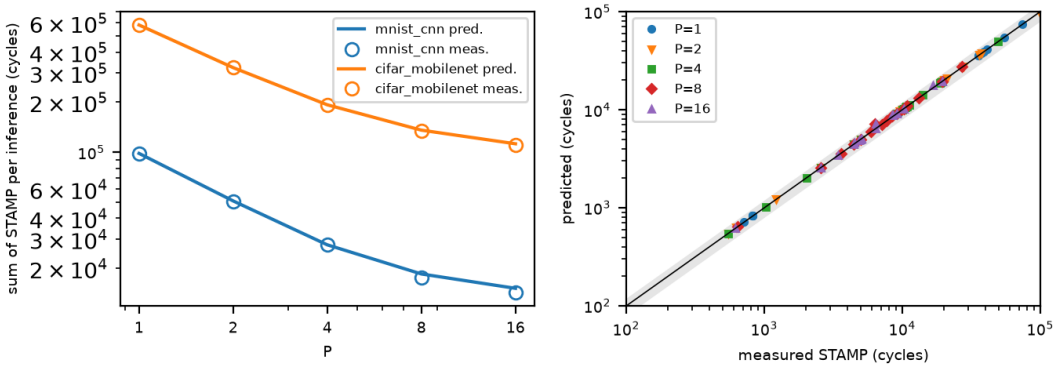



Fig. 2. Frozen prediction (tag `m1`, `os` architecture) against measurement. Left: summed STAMP per inference versus $P$ (line: prediction, circles: measurement). Right: every scored command, predicted versus measured cycles; the band is $\pm 20\,\%$.

TABLE V
PER-COMMAND CYCLES AT $P = 8$ (PRED. FROM TAG `m1`; MEAS. = MEDIAN STAMP OVER THE TIMED IMAGES).

| net | op | kind | K | pred. | meas. |
|---|---|---|---|---|---|
| mnist | 1 | conv | 3x3 | 7,144 | 6,356 |
| mnist | 3 | conv | 3x3 | 7,756 | 7,715 |
| mnist | 7 | fc | 1x1 | 3,596 | 3,594 |
| cifar | 0 | conv | 3x3 | 27,724 | 27,029 |
| cifar | 1 | dw | 3x3 | 10,236 | 10,203 |
| cifar | 2 | conv | 1x1 | 8,712 | 8,452 |
| cifar | 3 | dw | 3x3 | 19,436 | 19,376 |
| cifar | 4 | conv | 1x1 | 10,888 | 10,883 |
| cifar | 5 | dw | 3x3 | 5,012 | 4,995 |
| cifar | 6 | conv | 1x1 | 5,928 | 5,923 |
| cifar | 7 | dw | 3x3 | 9,204 | 9,176 |
| cifar | 8 | conv | 1x1 | 10,536 | 10,469 |
| cifar | 9 | dw | 3x3 | 2,580 | 2,571 |
| cifar | 10 | conv | 1x1 | 7,096 | 7,077 |
| cifar | 11 | dw | 3x3 | 4,436 | 4,424 |
| cifar | 12 | conv | 1x1 | 13,240 | 13,237 |
| cifar | 15 | fc | 1x1 | 644 | 642 |

law. The model has biases documented before `m1`; the setup bias proved larger than documented wherever a pending write is granted in the setup cycle (above), which is the kind of finding a frozen model makes visible and a fitted model would have absorbed. The baseline is a size-optimized multi-cycle core [21], so the speedups are relative to a deliberately small CPU and are not comparable with speedups over pipelined cores; they are reported to show the claims were met, not as a result. Only two networks are evaluated, both from one compiler, and one person wrote both the model and the RTL, so the independence between them rests on the git history rather than on separate authors; until the `m1` tag is timestamped outside the author's control, a reader has to trust that history.

## IX. Conclusion

We built a small int8 accelerator and its backend and scored an analytic cycle model that was frozen before the RTL. The model held to within 12.4 % per command and -0.9 % and 0.4 % per network, the resource budget was missed on LUTs, and every command that exceeded its registered bias has a mechanism identified post hoc. The outcomes of the pre-registered claims are those in the tables above; they update from `results/` through `experiments/make_macros.py`.

## Acknowledgment

AI assistance (Anthropic's Claude) was used to draft code and text; the author takes responsibility for all content.

## References

[1] H. Genc, S. Kim, A. Amid, A. Haj-Ali, V. Iyer, P. Prakash, J. Zhao, D. Grubb, H. Liew, H. Mao, A. Ou, C. Schmidt, S. Steffl, J. Wright, I. Stoica, J. Ragan-Kelley, K. Asanovic, B. Nikolic, and Y. S. Shao, "Gemmini: Enabling systematic deep-learning architecture evaluation via full-stack integration," in *58th ACM/IEEE Design Automation Conference (DAC)*, 2021, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1911.09925

[2] S. Prakash, T. Callahan, J. Bushagour, C. Banbury, A. V. Green, P. Warden, T. Ansell, and V. Janapa Reddi, "CFU Playground: Full-stack open-source framework for tiny machine learning (tinyML) acceleration on FPGAs," in *IEEE International Symposium on Performance Analysis of Systems and Software (ISPASS)*, 2023, pp. 157–167, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/2201.01863

[3] T. Moreau, T. Chen, L. Vega, J. Roesch, E. Yan, L. Zheng, J. Fromm, Z. Jiang, L. Ceze, C. Guestrin, and A. Krishnamurthy, "A hardware-software blueprint for flexible deep learning specialization," 2018, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1807.04188

[4] B. Jacob, S. Kligys, B. Chen, M. Zhu, M. Tang, A. Howard, H. Adam, and D. Kalenichenko, "Quantization and training of neural networks for efficient integer-arithmetic-only inference," 2017, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1712.05877

[5] B. Jacob *et al.*, "gemmlowp: fixedpoint/fixedpoint.h (saturatingrounding-doublinghighmul, roundingdividebypot)," GitHub, google/gemmlowp, accessed 2026-10-10. [Online]. Available: https://raw.githubusercontent.com/google/gemmlowp/master/fixedpoint/fixedpoint.h

[6] Google LiteRT (TensorFlow Lite) team, "Litert 8-bit quantization specification," Google AI Edge documentation, accessed 2026-10-10. [Online]. Available: https://developers.google.com/edge/litert/performance/quantization_spec

[7] R. David, J. Duke, A. Jain, V. Janapa Reddi, N. Jeffries, J. Li, N. Kreeger, I. Nappier, M. Natraj, S. Regev, R. Rhodes, T. Wang, and P. Warden, "TensorFlow Lite Micro: Embedded machine learning on TinyML systems," in *Proceedings of the 4th MLSys Conference*, 2021, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/2010.08678

[8] L. Lai, N. Suda, and V. Chandra, "CMSIS-NN: Efficient neural network kernels for Arm Cortex-M CPUs," 2018, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1801.06601

[9] J. Lin, W.-M. Chen, Y. Lin, J. Cohn, C. Gan, and S. Han, "MCUNet: Tiny deep learning on IoT devices," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/2007.10319

[10] Anirudh, "mcugrad: from-scratch deep learning framework and int8 C99 compiler for microcontrollers," Git submodule third_party/mcugrad at commit e9f6283, accessed 2026-10-10. [Online]. Available: https://github.com/anirudh12032008/mcugrad.git

[11] Y.-H. Chen, J. Emer, and V. Sze, "Eyeriss: A spatial architecture for energy-efficient dataflow for convolutional neural networks," in *43rd ACM/IEEE International Symposium on Computer Architecture (ISCA)*, 2016, accessed 2026-10-10. [Online]. Available: https://dspace.mit.edu/handle/1721.1/102369

[12] N. P. Jouppi, C. Young, N. Patil, D. Patterson *et al.*, "In-datacenter performance analysis of a tensor processing unit," in *44th International Symposium on Computer Architecture (ISCA)*, 2017, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1704.04760

[13] NVIDIA, "NVDLA hardware architectural specification: Unit description (nvdlav1)," nvdla.org, accessed 2026-10-10. [Online]. Available: http://nvdla.org/hw/v1/ias/unit_description.html

[14] Y. Umuroglu, N. J. Fraser, G. Gambardella, M. Blott, P. Leong, M. Jahre, and K. Vissers, "FINN: A framework for fast, scalable binarized neural network inference," in *25th International Symposium on Field-Programmable Gate Arrays (FPGA)*, 2017, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1612.07119

[15] A. Garofalo, M. Rusci, F. Conti, D. Rossi, and L. Benini, "PULP-NN: Accelerating quantized neural networks on parallel ultra-low-power RISC-V processors," 2019, accessed 2026-10-10. [Online]. Available: https://arxiv.org/abs/1908.11263

[16] A. Garofalo, G. Tagliavini, F. Conti, L. Benini, and D. Rossi, "XpulpNN: Enabling energy efficient and flexible inference of quantized neural network on RISC-V based IoT end nodes," 2020, accessed 2026-10-10. Preprint; the arXiv page lists no venue. [Online]. Available: https://arxiv.org/abs/2011.14325

[17] M. Yildirim and O. Ozturk, "RISC-V based TinyML accelerator for depthwise separable convolutions in edge AI," 2025, accessed 2026-10-10. Preprint, no peer-reviewed venue listed. [Online]. Available: https://arxiv.org/abs/2511.21232

[18] C. Zhang, P. Li, G. Sun, Y. Guan, B. Xiao, and J. Cong, "Optimizing FPGA-based accelerator design for deep convolutional neural networks," in *ACM/SIGDA International Symposium on Field-Programmable Gate Arrays (FPGA)*, 2015, accessed 2026-10-10. [Online]. Available: https://www.microsoft.com/en-us/research/?p=595783

[19] M. Alwani, H. Chen, M. Ferdman, and P. Milder, "Fused-layer CNN accelerators," in *49th Annual IEEE/ACM International Symposium on Microarchitecture (MICRO)*, 2016, accessed 2026-10-10. [Online]. Available: https://compas.cs.stonybrook.edu/~mferdman/downloads.php/MICRO16_Fused_Layer_CNN_Accelerators.pdf

[20] S. Williams, A. Waterman, and D. A. Patterson, "Roofline: An insightful visual performance model for multicore architectures," 2009, accessed 2026-10-10. [Online]. Available: https://zenodo.org/record/1236156

[21] C. Wolf, "PicoRV32: a size-optimized RISC-V cpu, readme," GitHub, YosysHQ/picorv32, accessed 2026-10-10. [Online]. Available: https://raw.githubusercontent.com/YosysHQ/picorv32/main/README.md

[22] Lattice Semiconductor, "ECP5 / ECP5-5G product page and family table," latticesemi.com, accessed 2026-10-10. [Online]. Available: https://www.latticesemi.com/en/Products/FPGAandCPLD/ECP5

[23] Anirudh, "corv32: five-stage RV32IM core and SoC (design docs, bus protocol)," Unpublished repository; vendored subset at third_party/corv32 of the mcuaccel repository, commit e1d0bc0, accessed 2026-10-10.